

Erica: AI Tutor

You: Can you explain the variational lower bound and how it relates to Jensen's inequality?

Erica:

Based on our course material, let's break down the concepts step-by-step starting from the prerequisites.

Prerequisite: Jensen's Inequality

Jensen's Inequality is a fundamental concept in convex analysis. It states that for a convex function f , and any random variable X :

$$f(E[X]) \leq E[f(X)]$$

Here, $E[X]$ represents the expected value of X .

Prerequisite: Evidence Lower Bound (ELBO)

The Evidence Lower Bound (ELBO) is a lower bound on the log marginal likelihood. In variational inference, we approximate the true posterior distribution $p(z|x)$ with an approximating distribution $q(z|x)$. The ELBO is derived using Jensen's inequality and provides us with a way to optimize our approximation.

Ask Erica something...

Send

Erica: AI Tutor

Variational Lower Bound (ELBO)

The Evidence Lower Bound (ELBO) can be written as:

$$\text{ELBO} = \mathbb{E}_{q(z|x)}[\log p(x, z)] - D_{\text{KL}}(q(z|x) \parallel p(z|x))$$

Where:

- $q(z|x)$ is the approximating distribution of the posterior.
- $p(z|x)$ is the true posterior distribution we are trying to approximate.
- $D_{\text{KL}}(* \parallel *)$ denotes the Kullback-Leibler divergence, which measures the difference between two distributions.

How ELBO Relates to Jensen's Inequality

To understand how ELBO relates to Jensen's inequality, let's rewrite the ELBO in a form that allows us to apply Jensen's inequality:

1. Logarithm of Joint Distribution:

$$[\log p(x, z) = \log ((p(x, z))/(p(z|x)p(x)) * p(z|x) p(x))]$$

Ask Erica something...

Send

Erica: AI Tutor

To understand how ELBO relates to Jensen's inequality, let's rewrite the ELBO in a form that allows us to apply Jensen's inequality:

1. Logarithm of Joint Distribution:

$$[\log p(x, z) = \log ((p(x, z))/(p(z|x)p(x)) * p(z|x) p(x))]$$

2. Apply Jensen's Inequality to the Expectation:

Since log is a concave function (the negative of a convex function), we can apply Jensen's inequality in its general form for concave functions:

$$E_{q(z|x)}[\log p(x, z)] \leq \log E_{q(z|x)}[p(x, z)]$$

3. Rewrite the Expectation:

We know that:

$$[E_{q(z|x)}[p(x, z)] = \int q(z|x) p(x|z) dz]$$

4. Combine Terms to Form ELBO:

By combining terms and rearranging, we get the Evidence Lower Bound (ELBO):

Ask Erica something...

Send

Erica: AI Tutor

$$[E_{q(z|x)}[p(x, z)] = \int q(z|x) p(x|z) dz]$$

4. Combine Terms to Form ELBO:

By combining terms and rearranging, we get the Evidence Lower Bound (ELBO):

$$[ELBO = E_{q(z|x)}[\log p(x, z)] - D_{KL}(q(z|x) \parallel p(z|x))]$$

The first term $E_{q(z|x)}[\log p(x, z)]$ is a lower bound on the log marginal likelihood $\log p(x)$, and we subtract the KL divergence to ensure that our approximation $q(z|x)$ does not deviate too much from the true posterior $p(z|x)$.

Conclusion

The Evidence Lower Bound (ELBO) provides us with a way to optimize our variational approximation by maximizing this lower bound. This is achieved using techniques like Variational Autoencoders, where we aim to minimize the negative ELBO, thus optimizing both the reconstruction loss and the KL divergence term.

In summary, Jensen's Inequality forms the basis for establishing the Evidence Lower Bound (ELBO), which in turn allows us to approximate complex posterior distributions effectively.

Ask Erica something...

Send

=====

127.0.0.1 - - [05/Dec/2025 20:42:13] "POST /ask HTTP/1.1" 200 -

Inside answer def

retrieve relevant thingy

Retrieved KG Nodes:

["Jensen's Inequality", 'Variational Inference', 'Variational Autoencoder Loss', 'Evidence Lower Bound (ELBO)', 'Empirical Risk Minimization']

=====

Saved KG subgraph visualization at static/query_graph.png

built prompt!

Prompt Sent to Model:

You are Erica, the AI tutor for Prof. Pantelis Monogioudis's "Introduction to AI" course.

RULES:

1. First use the course material below if relevant
2. If not enough in course material, add basic general knowledge
3. Always show any mathematical formulas or notation that exist for the topic
4. Keep explanations simple and clear
5. When explaining a concept, FIRST explain its prerequisites and show their mathematical formulas
6. Build up understanding step-by-step from foundational concepts to the current concept

WHEN ANSWERING:

- Start with: "Based on our course material"
- If adding info: "In general AI"
- Always show formulas: If topic has math, display it with $like this$
- For prerequisites: Explain each prerequisite concept with its own mathematical notation before connecting them to the main concept

COURSE MATERIAL:

Concept: Jensen's Inequality

Concept: Jensen's Inequality

Aliases: convexity inequality

Definition: A mathematical property stating that for a convex function, the function of an expectation is less than or equal to the expectation of the function. It is the core inequality used to derive variational bounds.

Concept: Variational Inference

Aliases: VI, variational methods

Definition: A technique that approximates intractable posterior distributions with simpler distributions by optimizing a lower bound, typically the ELBO.

Concept: Variational Autoencoder Loss

Aliases: VAE loss, ELBO loss

Definition: The objective function in VAEs consisting of a reconstruction loss and a KL divergence term, derived from maximizing the evidence lower bound (ELBO).

Concept: Evidence Lower Bound (ELBO)

Aliases: variational lower bound, ELBO

Definition: A lower bound on the log marginal likelihood derived using Jensen's inequality. Maximizing the ELBO allows variational models like VAEs to approximate the true posterior distribution.

Concept: Empirical Risk Minimization

Aliases: ERM

Definition: A principle in machine learning that involves finding a hypothesis with minimal expected error on the training data. It balances bias and variance to achieve good generalization performance.

QUESTION:

Can you explain the variational lower bound and how it relates to Jensen's inequality?

=====

127.0.0.1 - - [05/Dec/2025 20:45:38] "POST /ask HTTP/1.1" 200 -